

Where Campaigns Matter: Evidence from Labour’s 2024 General Election Campaign

Petra Laura Oreskovic*

Abstract

Political parties devote substantial resources to campaigning, yet we know little about the electoral effects of the campaigns they actually deploy. An identical campaign may also affect voters differently depending on their political preferences and the strategic choices they face in their constituency. I study the UK Labour Party’s 2024 general election campaign using internal data on constituency-level resource allocation. Labour’s allocation rule generates two discontinuities that allow me to estimate the causal effects of mobilizing its full set of campaign resources in both swing and stronghold seats. The campaign increased Labour’s victory margin by 6.8 percentage points in strongholds but had no comparable effect in swing seats. Individual-level data indicate that left-leaning voters consolidate behind Labour in both settings. The right-wing electoral bloc fragments in strongholds as voters shift toward the far right, but consolidates behind the Conservatives in swing seats. The electoral effect of campaigning is therefore endogenous to the competitiveness of a seat, reflecting not only support for the campaigning party but also how its opponents’ supporters respond.

1 Introduction

Campaigns are a central part of the democratic process, shaping who is represented, which parties govern, and ultimately which policies are pursued. Parties devote enormous resources to this effort. In the 2024 UK general election alone, parties and campaigners reported £94.5 million in spending (Electoral Commission, 2025), while spending in the 2024 US elections reached approximately \$15 billion (OpenSecrets, 2024). Despite the scale of modern campaigning, there remains substantial uncertainty about how much campaigns change electoral outcomes, whether their effects vary with the competitiveness of a race, and whose behavior they change.

*University of California, Berkeley. petra.laura@berkeley.edu. I am especially grateful to my committee, Francesco Trebbi, Federico Finan, and Cailin Slaterry, for their guidance throughout this project. I also thank Matilde Bombardini, Jake Grumbach, Gabriel Lenz, and Yotam Margalit for helpful comments. The latest version of this paper is available at <https://petra-laura-oreskovic.eu/papers/where-campaigns-matter.pdf>.

What is the effect of a campaign as run by a political party trying to win an election? Parties combine canvassing, advertising and candidate visits at considerable scale and intensity, drawing on internal research to determine whom to target and where to allocate their resources. Isolating the causal effect of such a composite campaign is challenging because (i) reproducing its scale is rarely feasible in an experimental setting and (ii) randomizing treatment assignment generally removes the party’s strategic decision about where to allocate its resources. Even experiments conducted within real campaigns generally vary only one part of the bundle of activities that a party engages in.

Campaign effects are also unlikely to be uniform across electoral settings. Where campaigns are effective and which voters they influence can affect parties’ electoral strategies, including the policies they adopt to attract voters (Downs, 1957; Dixit and Londregan, 1996). How campaign effects vary with the competitiveness of a district is particularly consequential because the value of an additional vote is higher in closer races. The relative number of a party’s base and swing voters and the importance of tactical voting also change with electoral competitiveness, potentially affecting both the magnitude of campaign effects and which voters drive them.

I study these questions in the context of the Labour Party’s campaign in the 2024 UK general election, which provides a rare opportunity to identify the causal effects of the full range of resources a major political party mobilizes to win an election. Specifically, I estimate the effect of Labour’s decision to prioritize certain constituencies for canvassing, digital and physical advertising, assigned organizers, and high-profile visits. Labour spent £30 million over the regulated campaign period, including almost £12 million on digital advertising (Electoral Commission, 2025), while its activists spoke to more than 3 million people (Black, 2024). Labour designed the campaign and decided where to concentrate its resources using constituency-level electoral predictions from its internal multilevel regression and post-stratification (MRP) model. This setting allows me to study the aggregate effect of the campaign the party itself chose to run, rather than a campaign element randomized by researchers.

Labour’s allocation of resources across constituencies also allows me to study whether the effectiveness of its campaign varies with the competitiveness of a seat. Constituencies were prioritized based on Labour’s predicted margin over its best-performing competitor, an output of the party’s internal MRP model. When the model predicted a tie, Labour sharply increased campaign prioritization. Constituencies remained prioritized until Labour’s predicted lead reached 15 percentage points, at which point prioritization fell sharply. I refer to these as the swing and stronghold thresholds, respectively. I exploit both discontinuities in treatment assignment in a fuzzy regression discontinuity design to estimate the causal

effect of campaign prioritization at each threshold. The two thresholds capture the same campaign in substantively different electoral settings, with potentially different electorates and patterns of competition among parties. This allows me to compare campaign effects across swing and stronghold constituencies while maintaining the same research design and treatment.

I combine the internal campaign data with official election results and the British Election Study panel, which surveys approximately 30,000 respondents per wave across multiple waves before and after the election. This allows me to study both constituency-level electoral outcomes and the individual-level responses that produce them. Combining the panel with the regression discontinuity design, I examine which voters change their behavior in response to the campaign, the parties they move between, and whether their political attitudes change.

The effects of the campaign differ markedly across the two allocation thresholds. At the stronghold threshold, prioritization increases Labour's victory margin by 6.8 percentage points. Conservative vote share falls by 6.0 percentage points and the Conservative victory margin by 8.2 points, while Labour's own vote share rises by 2.3 points but is imprecisely estimated. At the swing threshold, by contrast, I find no corresponding increase in Labour's vote share or victory margin. Even within the same election and under the same party-wide strategy, the electoral effect of concentrating campaign resources in a seat differs sharply between more and less competitive seats.

Why does the same campaign produce such different returns across the two thresholds? The contrast seems to come primarily from the campaign's effects on support for Labour's opponents. This is clearest within Labour–Conservative contests, where either of the two parties could win around the swing threshold, while the Conservatives have little prospect of winning around the stronghold threshold. Labour's estimated vote-share effect is positive, though imprecise, at both thresholds. In contrast, campaigning reduces Conservative support where Labour is comfortably ahead but increases it around the swing threshold. This reversal suggests that the electoral viability of the Conservatives determines how voters respond to Labour's campaign.

The British Election Study provides evidence that this pattern may reflect tactical coordination within electoral blocs. In Labour–Conservative contests, voters who supported the Liberal Democrats or Greens in 2019 consolidate behind Labour at both thresholds. Among 2019 Conservative voters, Labour's campaign increases support for Reform at the stronghold threshold, fragmenting the right-wing bloc. At the swing threshold, the campaign instead increases their support for the Conservatives. Progressive consolidation behind Labour is offset by consolidation on the right, substantially reducing the net electoral gain from the campaign in close races. The campaign appears to activate tactical behavior, consistent

with the aggregate patterns of tactical voting within electoral blocs documented by [Miori and Green \(2025b\)](#). The individual-level evidence also shows that the aggregate decline in Conservative support and increase in Labour support do not reflect direct Conservative-to-Labour persuasion, despite Labour’s explicit efforts to win former Conservative voters.

I test several alternative explanations for the difference in the campaign’s effects between the two thresholds. One possibility is that the electorates around the two thresholds differ in the share of voters predisposed to respond to Labour campaigning, so that differences in voter composition alone generate the contrasting effects. I find that differences in the distribution of political preferences across the two thresholds do not account for the difference in campaign effects. Another possibility is that Labour faces a stronger countervailing Conservative campaign in the closest contests. I construct a constituency-level measure of the Conservative Party’s central campaign effort using Meta advertising data. Because Meta does not disclose the geographic targets of individual ads, I extract place names from Conservative ad text and graphics, geocode those locations to constituencies, and combine them with disclosed ad spending to infer where the central party concentrated its advertising. Controlling for this measure of Conservative campaign activity makes little material difference to the results. I also find that differences in access to social networks of politically like-minded voters do not explain the contrasting effects of the campaign.

The campaign also changes voters’ assessments of the incumbent Conservative government. The government was already deeply unpopular, with an approval rating of just 10.4 percent before the election ([Miori and Green, 2025a](#)). At the stronghold threshold, the campaign produces more negative views of the Conservatives, the economy, and the government, and lower expectations that the Conservatives will achieve major policy goals. Evaluations of Labour itself do not improve. The campaign therefore appears to reinforce existing dissatisfaction with the incumbent rather than improve perceptions of Labour.

Finally, I find that the consequences of campaign prioritization extend beyond election day. Labour MPs in prioritized constituencies around the stronghold threshold were substantially more likely to call for Prime Minister Keir Starmer to resign during the 2026 leadership crisis. Starmer ultimately left office. This suggests that campaign-induced changes in a constituency may subsequently affect the political incentives facing its elected representative.

The paper contributes first to the large causal literature on campaign effects. Earlier experimental work on partisan campaign effects largely measured attitudes and electoral behavior through surveys ([Ansolabehere et al., 1994](#); [Ansolabehere and Iyengar, 1995](#); [Arceneaux, 2007](#); [Kalla and Broockman, 2018](#)), while field experiments used administrative voting records to establish the effects of canvassing and other forms of voter contact on turnout ([Gerber and Green, 2000](#); [Green and Gerber, 2004](#); [Nickerson, 2008](#); [Arceneaux and Nicker-](#)

son, 2009). More recent field experiments collaborate with political campaigns to randomize particular components of campaign activity. Pons (2018) randomizes door-to-door canvassing across precincts as part of a nationwide presidential campaign, while Kendall et al. (2015) and Cruz et al. (2026) vary the messages delivered by the campaign, from valence versus ideology to policy-focused versus emotional appeals. Canvassing and particular campaign messages are found to produce sizeable effects on vote shares in these settings. A related experimental literature randomizes television advertising (Gerber et al., 2011), online advertising (Hager, 2019; Coppock et al., 2022; Aggarwal et al., 2023; Enríquez et al., 2024), and street-level political advertising (Esteban-Casanelles, 2024), with effects on electoral behavior ranging from small or imprecise to sizeable.

Even when embedded in large campaigns and at high intensity, however, these experiments generally isolate particular campaign activities. Their effects need not sum to the effect of the composite campaign that parties deploy, which is central for understanding how campaigning ultimately changes electoral outcomes. Furthermore, random assignment also replaces the party’s decision about where to deploy that activity, removing from treatment assignment the information and strategic considerations that ordinarily shape campaign resource allocation.

Some studies preserve more of campaigns’ own choices. Alvarez et al. (2010) allow campaign discretion over voter contact, while Enos and Fowler (2018) study the aggregate deployment of presidential ground campaigning, identifying effects on turnout by comparing targeted and untargeted states within media markets. Spenkuch and Toniatti (2018) obtain cleaner quasi-experimental variation in presidential television advertising from FCC-induced differences in exposure across neighboring counties, while Larreguy et al. (2018) exploit analogous cross-state media spillovers for radio advertising in Mexico. These designs preserve more of parties’ actual campaign choices, but their causal variation again comes from exposure to individual campaign activities. In my setting, prioritization mobilizes the full set of campaign resources, and the resulting LATEs are identified at the substantively relevant margins where the party itself is deciding whether to allocate additional resources. This identifies the effect of a campaign as actually run by a political party trying to win an election.

Second, the paper helps explain why campaign effects vary across electoral settings. Estimates range from small or null persuasive effects (Kalla and Broockman, 2018) to large and electorally consequential effects of particular campaign activities (Kendall et al., 2015; Pons, 2018; Enríquez et al., 2024). Existing work has examined several sources of this variation. Hewitt et al. (2024) study differences in campaign content, while Broockman and Kalla (2023) emphasize voters’ prior information about candidates and the informational

content of campaign messages. However, we might particularly care about how campaign effects vary with electoral competitiveness because votes gained through campaigning can more directly determine electoral outcomes in close races. Competitiveness has received some attention as a source of heterogeneity in campaign effects, but typically as one of several factors examined within studies of campaign advertising. [Coppock et al. \(2020\)](#) examine differences across electoral contexts, including battleground and non-battleground states, but study individual advertisements and survey responses. [Larreguy et al. \(2018\)](#) focus on the informational role of radio advertising about non-dominant parties and find larger effects where voters are less informed, including in less competitive precincts.

The two-threshold design provides a particularly direct test of whether the same composite campaign has different effects in more and less competitive races. The results help explain the wide range of estimates of campaign effects by showing that these effects are endogenous to the competitive environment in which campaigning takes place. I find that the same campaign changes voters' choices at both thresholds but produces sharply different net electoral returns as competitiveness changes voters' incentives to vote tactically. This suggests that campaigns may be least effective at generating net electoral gains in the close races where they are most consequential.

Third, I can provide a particularly detailed account of why the campaign's effects differ across the two thresholds. A large experimental literature distinguishes between mobilization and persuasion, studying which voters are induced to turn out by campaign contact ([Gerber and Green, 2000](#); [Arceneaux and Nickerson, 2009](#); [Cantoni and Pons, 2021](#)) and whether campaigns persuade voters to change their vote choice ([Arceneaux, 2007](#); [Kendall et al., 2015](#); [Bailey et al., 2016](#); [Kalla and Broockman, 2018](#)). Less attention has been paid to how campaigns affect the way voters translate their preferences into votes for particular parties through tactical voting. [Núñez \(2025\)](#) is an important exception, finding that party contact in British elections can encourage voters to remain with their preferred party or defect to a more viable alternative.

I combine quasi-random variation in campaign prioritization with an unusually rich individual panel, allowing me to identify the movements across parties that generate the aggregate electoral effects. I can identify which groups change their votes and where those votes go, revealing distinct patterns of consolidation across groups in a multiparty context. The distinction between direct persuasion and these patterns of reallocation matters substantively. The aggregate decline in Conservative support might otherwise be interpreted as direct Conservative-to-Labour persuasion, whereas the panel data instead point primarily to reallocation within ideological blocs.

Finally, the paper contributes to the large literature on electoral accountability and politi-

cians’ behavior in office. Much of the causal evidence on electoral incentives exploits term limits, comparing politicians who can seek reelection with those who cannot. Reelection incentives have been shown to affect economic policy choices (Besley and Case, 1995; Alt et al., 2011; List and Sturm, 2006), corruption (Ferraz and Finan, 2011), and legislators’ effort in office (Fouirnaies and Hall, 2022). Work on British MPs similarly finds that electoral vulnerability is associated with greater parliamentary activity (Bowler, 2010; Kellermann, 2013, 2016). In contrast, I provide causal evidence that a party’s decision to prioritize a constituency during an election can affect the subsequent behavior of its elected representative. Campaign prioritization changes the composition of electoral support in a constituency, including by drawing former Liberal Democrat and Green voters toward Labour. MPs elected from prioritized constituencies subsequently become substantially more likely to call for their own prime minister to resign during a leadership crisis that culminated in his departure from office. The consequences of campaigning can therefore extend beyond voters and election results to the political behavior of elected representatives.

2 Setting

2.1 2024 election

The 2024 election was announced by then Prime Minister Rishi Sunak on May 22, 2024. The announcement came as a surprise to the public and to Labour alike, as the election was widely expected to be held in the fall (Ford et al., 2025, p. 230). The election was called for July 4, 2024. The sudden announcement meant the campaign would run for just 6 weeks, and Labour scaled up its campaigning effort considerably in response (Ford et al., 2025, p. 293). Labour entered the campaign with a large polling lead, but Labour’s campaign leadership nonetheless expected the lead to narrow as election day approached.

Labour ultimately won 33.7% of the vote and 411 of the 650 seats in the House of Commons, its third-highest seat total in history despite a relatively low national vote share. The Conservatives won 121 seats on 23.7% of the vote, their worst result ever. The Liberal Democrats won 72 seats on 12.2%, Reform UK won five seats on 14.3%, and the Green Party won four seats on 6.7%. Four independent candidates in areas with large Muslim populations also defeated Labour candidates, standing in protest of the party’s position on the war in Gaza (Cracknell and Baker, 2024).

2.2 Labour’s 2024 targeting rule

Elections to the Parliament of the United Kingdom are governed by a first-past-the-post electoral system, meaning that the party that gains a plurality of votes in a constituency wins that seat in parliament. As campaign resources are scarce, parties must decide which constituencies to prioritize and allocate campaign funds and effort to. Identifying those seats is challenging, as even large national polls lack sufficient constituency-level samples, while estimates based on past results and an assumption of a uniform national swing may fail to capture changes in which seats are competitive.

To obtain reliable constituency-level predictions, contemporary parties in FPTP systems use Multilevel Regression and Post-stratification (MRP) models (Ford et al., 2025, p. 297). These models combine large national poll samples with administrative data and use multilevel models to partially pool estimates across similar constituencies where local samples are sparse. Predictions for different demographic types are then weighted by their share of each constituency’s electorate.

Labour developed its own internal MRP model ahead of the 2024 election. I worked on the party’s internal data science team during this period and through the first year of the Starmer government. Both my knowledge of the targeting rule and the internal prioritisation data used in this paper come from that experience.

The Labour Party tied its campaign targeting directly to the output of its internal MRP model. The model predicted Labour’s margin over the best performing competitor in each constituency, and the party targeted those where the predicted margin fell between 0 and 15 percentage points. Prioritisation increased sharply when Labour’s predicted margin crossed zero and fell sharply once its predicted lead reached 15 percentage points. I refer to these thresholds as swing and stronghold, respectively. These thresholds are visible in the prioritisation data as sharp changes in the probability of treatment, which I document in Section 4.

Assignment to the campaign treatment was not exclusively determined by the MRP output. Because the MRP partially pools estimates across similar constituencies where local samples are sparse, it could understate circumstances specific to an individual seat. This was particularly consequential where Labour candidates faced strong pro-Gaza independent challengers, leading the party to pay particular attention to those seats and exempt them from MRP-recommended treatment status. The campaign also considered whether a constituency had previously returned a Labour MP, whether it was rural, and the strength and recent electoral performance of the local Labour organization, including Labour’s share of local councillors and its win rate in recent local elections. The MRP-predicted margin was nonetheless the primary determinant of prioritization.

Treatment in this design starts on 22 May 2024, when Sunak called the election and trig-

gered the short campaign. This is the point at which the party made its resource allocation decision. A round of adjustments on 19 June further restricted the prioritization to previously prioritized seats close to the swing threshold. I do not use this round as treatment. By then, campaign effort had already been running for four weeks, so it may reflect the effect of the campaign at that point.

Figure 1 provides an initial indication that prioritization may have affected electoral outcomes. Prioritized constituencies lie predominantly above the fitted relationship between Labour’s predicted and realized margins, suggesting that Labour performed better in these constituencies than its predicted margin alone would imply.




Figure 1: The relationship between Labour’s predicted margin and the realized victory margin.

2.3 Labour’s campaign and spending rules

The 2024 campaign recorded the highest election spending on record. Labour spent £30 million, the Conservatives £23 million, and £94.5 million in total was spent including spending by non-party organizations ([Electoral Commission, 2025](#)). The UK enforces spending limits on campaigning by parties and candidates separately. Parties could spend up to £54,010

per constituency contested, giving Labour a national spending limit of approximately £34 million. Party spending could be allocated geographically, while spending associated with a particular candidate was subject to a separate local limit of approximately £19,000 on average.¹ Labour’s total spending approached the national limit for parties. Its largest spending category was unsolicited materials, including leaflets and direct mail sent to voters. Labour also spent close to £12 million on digital advertising, with additional spending on message testing, polling, and analysis.

In addition to financial resources, parties can mobilize their members for direct voter contact. During the six-week campaign, 16% of Labour members took part in face-to-face canvassing or telephone campaigning (Ford et al., 2025). Labour activists spoke to 3.1 million voters (Black, 2024), made 300,000 phone calls, and delivered 70 million pieces of campaign literature (Ford et al., 2025).

The campaign manager, and future Chief of Staff to the Prime Minister, Morgan McSweeney, directed the campaign. The leadership decided to prioritize some seats for physical and digital advertising and, crucially given the overall spending limits, canvassing effort. This involved moving party members to seats where the central campaign deemed their door-knocking would generate the greatest gains for the party. The campaign included “throwing everything into such seats, ‘Keir visits, as much money as they needed, organisers’” (Ford et al., 2025, p. 294). Prioritization thus concentrated the full range of resources at the party’s disposal in these constituencies, including spending, advertising, organizers, leadership visits, and canvassers. The resulting treatment is a composite measure of the party’s targeted campaign at close to maximum intensity.

The central party took several steps to enforce this allocation. Ellie Reeves, later Chair of the Labour Party, personally checked in with parliamentary candidates to keep them campaigning in seats other than their own (Ford et al., 2025, p. 294). Local parties in non-target seats also reported having access to Contact Creator, the party’s centralised canvassing system, cut off.² Without Contact Creator, canvassers could not identify which addresses to target or record the information gathered at the doorstep, severely limiting their ability to conduct targeted canvassing.

The party also tested messages used in the campaign and identified key groups in the electorate it aimed to target. Labour’s electoral strategy focused heavily on constituencies previously held by the Conservatives, reflecting the geography of party support under the UK’s electoral system (Ford et al., 2025, p. 281). Within these seats, the campaign focused primarily on Conservative voters who might be open to switching to Labour (Ford et al.,

¹The limit varies based on the type of seat, with borough and county seats subject to different limits.

²As reported by the elections analyst Mark Pack; see also Ford et al. (2025, p. 294).

2025, p. 287). Some of these voters were working class, socially conservative, and supportive of Brexit (Ford et al., 2025, p. 281). However, such voters were sparse in some constituencies and the available individual-level data were imprecise. The pool of voters targeted through door-knocking was therefore broader and included voters from other groups.

2.4 The Conservative campaign

The Conservative campaign was defensive throughout almost the entire campaign. CCHQ stopped funding seats it did not already hold from early 2022, ending an earlier strategy built around expanding on the 2019 majority (Ford et al., 2025, p. 296). By the short campaign, CCHQ’s stated goal was to “get the base vote up to a point where well-performing MPs... would be able to win” (Ford et al., 2025, p. 296).

Of the £23 million spent in total, the party allocated £2.46 million to Meta and Google ads, against Labour’s £6.62 million, and concentrated its spending overwhelmingly on Meta (Ford et al., 2025, p. 325). Digital advertising included 70 distinct negative ads in the first three weeks of June, targeted at more than a thousand small localities and later at older voters in traditionally Conservative areas (Ford et al., 2025, pp. 330, 332). These data are used to measure Conservative party targeting and are described further in the Data section.

Beyond digital advertising, which CCHQ directed centrally, the strength of a seat’s ground campaign often depended on local resources, including members willing to knock on doors and funding for campaign staff (Ford et al., 2025, p. 296). The activist base was often elderly and more inclined to donate than canvass (Ford et al., 2025, p. 296). Only 12 per cent of Conservative members reported canvassing by phone or in person, the lowest share of any major party (Ford et al., 2025, p. 308). In previous elections, local campaigns could also draw on MPs and activists from safe Conservative seats through a system known as “Mutual Aid.” By 2024, however, almost no Conservative seat was considered safe enough for its MP to leave and campaign elsewhere (Ford et al., 2025, p. 302).

The party’s websites also carried volunteer pages for canvassing, Calling Connect for phone canvassing and Walking2Win for door-to-door canvassing. These provide additional information on the geographic distribution of Conservative canvassing.

I measure this allocation directly rather than relying on these survey and spending aggregates. Using digital advertising data and Labour’s own records of where Conservative canvassing was taking place, I measure the geographic distribution of Conservative campaign effort and use it in the analyses that follow. The available accounts and these measures show a defensive allocation of resources, with greater Conservative activity in areas of higher predicted Conservative support. There is no indication of a system comparable to Labour’s for

reallocating resources across constituencies and based on particular thresholds, and neither measure shows a visible discontinuity at Labour’s prioritization thresholds (see Sections 4.4.3 and 6.1).

3 Data

Data on Labour’s targeting decisions come from two internal party datasets that, to my knowledge, no outside researcher has had access to before.

The first dataset contains the MRP estimates. For each constituency, the MRP model produces a mean predicted vote share for each party, along with a standard deviation reflecting the uncertainty around that prediction. From these, the campaign constructs Labour’s margin over the best performing competitor in each constituency, and this margin is the running variable in the RD design. The second is a separate dataset recording the party’s actual resource allocation decision, i.e. whether a constituency was designated a priority seat for canvassing, digital and physical advertising, and campaign visits.

Prioritization was not a deterministic function of the MRP margin alone. As described in Section 2.2, I observe several constituency characteristics that entered the campaign’s allocation decisions, and I control for these in the main specifications. They include an indicator for a high Muslim population share, whether the constituency had previously returned a Labour MP, whether it is rural, Labour’s share of local councillors, the share of wards holding elections, and Labour’s win rate in recent local elections.

I measure the existing strength of the local party using the share of local councillors who are Labour, based on Open Council Data, and recent electoral performance using Labour’s win rate in local elections. Because only a subset of wards holds elections in a given year, I also control for the share of wards holding elections. The demographic and rural measures are constructed from Census data. I additionally control for whether the seat had already been prioritized before the short, high-intensity campaign began, and for notional 2019 Conservative and Liberal Democrat vote shares to account for differences in prior electoral support across constituencies.

I also use data from the British Election Study, waves 26 through 29 (May 2024 through July 2024, with wave 29 covering the post-election period and reporting actual vote choice as opposed to vote intention), with approximately 30,000 respondents per wave. I draw on items covering past vote choice, beliefs about the parties and the economy, political discussion, and self-reported strategic voting. Depending on the analysis, I use a wave’s full cross-sectional sample or subset to the panel of respondents present across specific waves. Some items are asked only post-election and are analysed on that wave alone. Some analyses

define subgroups from an earlier, pre-campaign wave and examine their outcomes later.

I link the MRP data, allocation data, BES respondents, constituency characteristics, and measures of Conservative campaign activity at the constituency level.

I use Meta advertising data to measure the Conservative Party’s central campaign strategy. National party spending reflects central strategy better than candidate-level spending does, which depends more heavily on local resources. National spending across larger parties and campaigners reached £69.3m in 2024, roughly three times the £23m spent by candidates across all parties combined. While most national spending is allocated to unsolicited mailers, whose geographic distribution I cannot observe, I assume that online spending follows similar geographic patterns.

Parties can target Meta ads down to the postcode level, but the platform does not publicly disclose where an ad was actually targeted, only what was spent on it. To recover the constituency-level geography of the Conservatives’ Facebook advertising, I scrape all ads paid for by the central party and apply OCR to their creative graphics, including videos, to extract place names. Many Conservative ads follow one of a small number of fixed templates in which only the named town varies. For example, one format states: “Don’t risk a Keir Starmer supermajority. Don’t let your community in [town] lose its voice.” I extract a named town from 9,424 of the 12,578 ads paid for by the Conservative Party. Most of the unmatched ads do not contain an identifiable place name. I then geocode the town name and assign it to a constituency, and combine this with each ad’s disclosed spend to construct Meta ad spending totals by constituency.

This is a novel measure of party effort. It is not universally applicable, as it depends on a party running ads in a formulaic style that includes identifiable place names. However, it improves on using contact rates reconstructed from survey responses to the British Election Study, the standard approach in [Ford et al. \(2025\)](#) and other work. Estimating campaign contact reliably at the constituency level would require precise estimates for each of roughly 630 constituencies, more than even the BES’s large sample can support.

I also use Labour’s own internal monitoring of Conservative campaign activity, based on records of postings under the Conservative Party’s Walking2Win, Ring2Win, and Knock and Drop Survey tools, recording session title, postcode, date, and constituency. These include 896 door-knocking sessions in the pre-election period. A posted session means a slot was scheduled, but it does not record the intensity of the effort.

For additional analyses, I rely on LabourList’s tracker of Labour MPs’ declared positions during the 2026 leadership crisis. As of 14 May 2026, 158 MPs had shown support for Starmer, 103 had called for his resignation, and 143 were unknown or undeclared. I use this to test a downstream consequence of prioritization, i.e. whether receiving campaign resources

made an MP more or less likely to back Starmer’s removal.

4 Empirical strategy

4.1 Sample construction

The initial sample consists of 632 constituencies in Great Britain. Labour does not contest elections in Northern Ireland, nor do the other parties considered in the analysis, and Labour’s MRP model therefore produces predictions only for England, Wales, and Scotland.

The main analysis is restricted to constituencies in which the pre-campaign MRP prediction identified Labour, the Conservatives, Reform, or the Liberal Democrats as the relevant competitive parties. In seats involving smaller parties such as Plaid Cymru or the Greens, partial pooling can obscure unusually strong constituency-specific support. A smaller party may therefore substantially outperform its MRP prediction in a particular seat, meaning that the constituency’s predicted position on the running variable differs considerably from its underlying electoral competitiveness. Such prediction error could otherwise be mistaken for an effect of the campaign.

I also exclude Scottish constituencies, where differences in the structure of electoral competition create a similar risk of attributing underlying competitive patterns to the campaign. Appendix Figure A1 plots Conservative vote share against Labour’s MRP-predicted margin over its best performing competitor and shows that Scottish constituencies cluster at a much lower level of Conservative support throughout the distribution of the running variable. In Scotland, the SNP is typically Labour’s best performing competitor rather than the Conservatives, so the running variable captures a different competitive dynamic. The main analysis therefore focuses on England and Wales and further restricts the sample to bandwidths around the two thresholds described below.

4.2 MRP-driven rule as a source of exogenous variation

Labour’s allocation of campaign resources provides two sources of discontinuous variation in campaign prioritization. As described in Section 2.2, constituencies were prioritized based on their predicted Labour margin over the best performing competitor in the party’s internal MRP model. Prioritization refers here to the classification in place at the start of the campaign, which determined the allocation of canvassing, advertising, and visibility resources across constituencies. Prioritization changes sharply at two predetermined thresholds. The swing threshold corresponds to a predicted tie between Labour and its best performing competitor, while the stronghold threshold corresponds to a predicted Labour lead of 15

percentage points. I exploit these discontinuities in a fuzzy regression discontinuity design to estimate separate local effects of campaign prioritization at the two thresholds on vote share, turnout, and a party’s victory margin.

A small internal team built and ran the MRP model with the sole objective of predictive accuracy. The team was institutionally separate from campaign decisions regarding individual constituencies and the model drew on national poll microdata and census data. No constituency party supplied or controlled those inputs, so they could not manipulate the inputs to the model nor precisely determine their position relative to either threshold. Furthermore, campaign resources are finite, and the party rigorously enforced the prioritization rather than spreading canvassers and advertising continuously across all 632 seats (Section 2.2).




Figure 2: Prioritization status as a function of the MRP-predicted lead.

Figure 2 plots prioritization status against the predicted margin over the best performing competitor for Labour. The dotted vertical lines mark the swing and stronghold cutoffs, while dots denote individual constituencies. Figure 2 shows clear discontinuities in prioritization status at both cutoffs.

Prioritization was not shaped exclusively by the MRP output. As described in Section 2.2, the campaign sometimes departed from the MRP rule in response to constituency-specific circumstances and other information it considered relevant to allocation.

The MRP margin followed a strict mathematical rule, but the campaign’s own judgement that occasionally overrode the MRP output reflected a mental model rather than an exact

formula. Formalizing that judgement into an explicit decision rule would be both difficult and inherently somewhat arbitrary. I instead control directly for the observable factors identified as relevant by the campaign, described in Section 3. I include an indicator for a high Muslim population share, whether the constituency had previously returned a Labour MP, rural status, Labour’s share of local councillors, the share of wards holding elections, and Labour’s win rate in recent local elections. I also control for whether the seat had already been prioritized before the short campaign began, to isolate the effect of prioritization during the main campaign, and for notional 2019 Conservative and Liberal Democrat vote shares.

The MRP margin was nonetheless the primary determinant of prioritization, and prioritization changes sharply at both thresholds. Section 5 reports the corresponding first-stage estimates.

4.3 Empirical specifications

The identifying assumption is that campaign prioritization is the only determinant of electoral outcomes that changes discontinuously at either threshold. Under this assumption, crossing either threshold provides quasi-random variation in campaign prioritization locally around the cutoff. I estimate the following within a single window spanning both cutoffs:

$$\text{FS: } D_i = f(X_i) + \gamma_1^{FS} Z_{1i} + \gamma_2^{FS} Z_{2i} + W_i' \beta^{FS} + \varepsilon_i \quad (1)$$

$$\text{RF: } Y_i = f(X_i) + \gamma_1^{RF} Z_{1i} + \gamma_2^{RF} Z_{2i} + W_i' \beta^{RF} + \varepsilon_i \quad (2)$$

Here X_i is the predicted margin, $Z_{1i} = \mathbf{1}\{X_i \geq 0\}$ and $Z_{2i} = \mathbf{1}\{X_i > 0.15\}$ mark which side of the swing and stronghold cutoffs a constituency sits on, and D_i is whether the seat was actually prioritized. Equation (1) estimates the change in the probability of prioritization at each cutoff. Equation (2) estimates the corresponding discontinuity in the electoral outcome Y_i .

W_i is the control vector from Section 3. It includes an indicator for a high Muslim population share, whether the constituency had previously returned a Labour MP, rural status, Labour’s share of local councillors, the share of wards holding elections, Labour’s win rate in recent local elections, and whether the seat was already prioritized before the short campaign. I additionally control for notional 2019 Conservative and Liberal Democrat vote shares.

$f(X_i)$ is the running-variable term, estimated in two versions: an interaction specification, which allows the slope of the running variable to differ across the three segments defined by the two cutoffs, and a quadratic specification. I use the interacted linear specification as my

main specification and report the quadratic version as a robustness check.

Following [Hahn et al. \(2001\)](#), I combine the first-stage and reduced-form estimates into a Wald ratio at each cutoff:

$$\tau_j = \gamma_j^{RF} / \gamma_j^{FS}, \quad j \in \{1, 2\} \quad (3)$$

τ_1 and τ_2 identify local average treatment effects for compliers at each cutoff.

I estimate the two cutoffs jointly rather than in separate regressions, which improves precision in a relatively small sample without imposing a common treatment effect across the two thresholds. I deliberately allow the discontinuity to differ at each cutoff, because the swing and stronghold thresholds represent substantively different competitive environments and a single pooled effect would mask precisely the heterogeneity the paper is designed to identify. This is also consistent with [Cattaneo et al. \(2016\)](#), who show that pooling across multiple cutoffs can obscure treatment-effect heterogeneity. I also estimate each cutoff separately using CCT-selected local bandwidths and report these estimates alongside the main specification in Section 5 and the Appendix.

For clarity, I report the main estimates at the $[-0.20, 0.35]$ bandwidth in the figures that follow. The full set of bandwidths is reported in Appendix Figures A2–A5.

4.3.1 Manipulation

Precise manipulation of the running variable is unlikely by construction. Constituencies could not manipulate their predicted margin and precisely sort themselves above or below either assignment threshold, as discussed in Section 4.2.

I test this formally using the density discontinuity test of [Cattaneo et al. \(2018\)](#). Figure 3 plots the density of the running variable around each cutoff. The test fails to reject continuity at both the swing and the stronghold threshold, providing no evidence of sorting around either cutoff.

4.3.2 Covariate balance test

I next test whether predetermined constituency characteristics change discontinuously at either threshold. Figure 4 reports standardized RD estimates for constituency demographics, previous electoral outcomes, Labour’s local organizational strength, and electorate size. The estimates are generally centered around zero at both thresholds, with no consistent pattern of imbalance across characteristics. The overall pattern is similar across the broadest and





Figure 3: Density discontinuity (`rddensity`) test.

narrowest fixed windows and the data-driven bandwidth selected by `rdrobust`, although estimates for some of the previous election results vary across specifications. At the swing threshold, the Liberal Democrat 2019 notional vote share is statistically significant at the 5 percent level in the main specification, while the Conservative 2019 notional vote share is statistically significant in the other two specifications. However, the estimated discontinuities are not stable across specifications and they both change sign depending on the bandwidth used. I control for both measures throughout the analysis. None of the characteristics is statistically significant at the stronghold threshold.




Figure 4: Covariate balance.

I also test balance jointly across the predetermined constituency characteristics. I handle January–May prioritization separately because it is an earlier campaign decision rather than a pre-determined characteristic of the constituency, but I continue to control for it throughout the analysis. I fail to reject joint balance at either the swing ($\chi^2(11) = 17.3$, $p = 0.101$) or stronghold threshold ($\chi^2(11) = 9.3$, $p = 0.597$), or across both thresholds jointly ($\chi^2(22) =$

27.3, $p = 0.202$). The joint tests thus do not indicate systematic imbalance around either cutoff.

Finally, I assess whether the differences in predetermined characteristics around the thresholds predict meaningful differences in the outcomes. For each outcome, I combine the characteristics used as controls in the main specification into a covariate index, weighting each characteristic by its association with that outcome. The estimated discontinuities in these indices are less than one percentage point for Labour and Conservative vote shares and victory margins, and none is statistically significant. I control for these predetermined characteristics throughout the analysis.

4.3.3 Exclusion restriction

For the Wald estimator to identify the causal effect of Labour’s campaign, crossing the prioritization threshold must affect electoral outcomes only through the resulting change in Labour campaign activity. This interpretation also requires that the increase in Labour campaigning does not induce a discontinuous response by competing parties or other actors. Other determinants of electoral outcomes should be continuous at the prioritization thresholds (Imbens and Lemieux, 2008; Lee and Lemieux, 2010). Otherwise, the Wald estimate would capture the combined effect of Labour’s additional campaign activity and the campaign response it generates, rather than the effect of Labour’s campaign alone. I cannot directly test this assumption, but I provide several pieces of evidence in its support.

While Labour’s internal MRP thresholds were not known to other parties, they could potentially observe greater Labour activity in a constituency and respond by reallocating their own resources. This possibility is particularly acute for the Conservatives, who were Labour’s principal competitor in most of the constituencies in my sample. They had the strongest electoral incentive to respond to an increase in Labour campaign activity and the greatest scope for such a response. For the remaining seats, comparable data on other parties’ local activity is not available, so this test cannot speak to a strategic response from non-Conservative competitors.

I examine this possibility using two measures of Conservative campaign activity, constituency-level Facebook advertising expenditure and door-knocking events advertised on local Conservative pages that were posted on or after May 22, when the Labour prioritization took place.

Figure 5, panels (a) and (b), plot both measures against Labour’s predicted margin. There is no visible discontinuity in either measure at either prioritization threshold, and the estimated differences at the cutoffs are statistically indistinguishable from zero. The results provide no evidence that the Conservatives adjusted their campaign strategy in response



Figure 5: Conservative actions as a function of the Labour lead. Panel (a): Facebook advertising expenditure. Panel (b): door-knocking events.

to Labour’s prioritization, either by increasing activity in constituencies Labour prioritized or shifting it to those it did not. These two measures cover the main online and in-person dimensions of local campaign activity. There is little reason to expect the Conservatives to have shifted their response specifically into channels not captured here, such as direct mail. This supports interpreting the Wald estimates as effects of Labour’s campaign rather than a discontinuous Conservative response, although responses through unobserved channels cannot be ruled out with complete certainty.

Another potential concern is that the media respond to Labour’s campaign. However, coverage arising from campaign events or candidate appearances, as well as advertising in local media, are part of Labour’s campaign itself. On the other hand, independent changes in editorial coverage are unlikely to occur at the prioritization thresholds, as media markets rarely correspond to individual constituencies and are typically regional or national.

A related possibility is that voters respond directly to Labour’s prioritization, by glean- ing from increased activity or directly learning from the campaign that a race has become competitive and adjusting their vote accordingly. Because this response runs through the campaign activity voters observe in the constituency, it is part of the effect of Labour’s campaign that this design captures.

4.3.4 Monotonicity

The Wald estimator identifies a local average treatment effect only if crossing the priori- zation threshold never reduces the campaign activity a constituency receives (Imbens and Angrist, 1994; Imbens and Lemieux, 2008), and vice versa at the threshold where the Labour campaign turns off. This rules out defiers, constituencies whose activity would have been higher had they not been prioritized. A defier would require local effort to fall by more than the additional staff and funding provided through prioritization, leaving total cam- paign activity lower than it would otherwise have been. This is implausible in my setting, as prioritization added staff and funding on top of constituencies’ existing campaigns rather than replacing the local activity, so local effort would have to collapse by more than the additional resources for total activity to fall.

5 Results

In this section, I first present the first-stage relationship between the running variable and treatment assignment. Second, I report reduced-form estimates of the effect of crossing the treatment assignment thresholds on Labour, Conservative, Liberal Democrat, and Re- form vote shares and victory margins at the 2024 general election, alongside results on

turnout. Third, I combine the two into Wald (fuzzy RD) estimates at my preferred bandwidth of $[-0.20, 0.35]$. Fourth, I explore heterogeneity by contest type, splitting the sample by whether the race is a Conservative–Labour contest. Fifth, I summarize a battery of robustness checks. Sixth, I present corroborating individual-level evidence from the British Election Study (BES). Seventh, I consider mechanisms alternative to strategic voting, including changes in other parties’ campaign behaviour and compositional shifts in the electorate. Finally, I examine downstream consequences of the campaign for MPs’ political behaviour, focusing on support for Keir Starmer’s resignation.

Table 1 reports the headline results in full: first-stage, reduced-form, and Wald estimates at my main specification, together with the corresponding first-stage F -statistic, for each outcome and threshold. The subsections below walk through each column in turn before turning to heterogeneity, robustness, and corroborating evidence.

Table 1: First-stage, reduced-form, and Wald estimates

	Labour vote share (1)	Conservative vote share (2)	Labour victory margin (3)	Conservative victory margin (4)
Panel A: Swing threshold				
First stage:	38.82*** (8.53)		F -statistic: 20.7	
Reduced form	−1.23 (1.12)	−0.20 (1.08)	−3.19* (1.78)	−4.13** (1.93)
Wald	−3.16 (3.08)	−0.52 (2.77)	−8.22 (5.17)	−10.65** (5.06)
Panel B: Stronghold threshold				
First stage:	−53.34*** (10.77)		F -statistic: 24.5	
Reduced form	−1.56 (1.02)	2.96** (1.21)	−3.91** (1.69)	4.35** (1.74)
Wald	2.92 (2.04)	−5.54** (2.57)	7.33* (3.75)	−8.15** (3.70)
N	374	374	374	374

Notes: Sample is all races where Labour, the Conservatives, the Liberal Democrats, or Reform is a top-two contender in Labour’s internal MRP prediction. Estimates in percentage points at the main specification $[-0.20, 0.35]$, interaction). HC1 standard errors in parentheses. The first stage does not vary by outcome; its F -statistic is the squared t -ratio on that threshold’s jump. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

5.1 First stage

Before turning to the effect of treatment on voting outcomes, I first establish that crossing the threshold reliably predicts treatment itself. Following specification (1) but without controls, crossing significantly predicts prioritization at both thresholds, with a coefficient of 0.393 at the swing threshold and -0.617 at the stronghold threshold. The R^2 is 0.47, indicating that the thresholds predict prioritization strongly, albeit not deterministically. Adding controls for exceptions to the targeting rule such as rurality, the number of councilors, and prior electoral performance narrows the standard error at both thresholds, from 0.101 to 0.086 at the swing threshold and from 0.113 to 0.107 at the stronghold threshold, and raises the R^2 to 0.60.

At my preferred bandwidth of $[-0.20, 0.35]$, crossing significantly predicts entering treatment at the swing threshold, with a coefficient of 0.386, and exiting treatment at the stronghold threshold, with a coefficient of -0.533 . The first-stage F-statistics are 20.7 at the swing threshold and 24.5 at the stronghold threshold, both above the rule-of-thumb F-statistic threshold of 10 (Staiger and Stock, 1997).

Figure 6 shows the discontinuity directly, plotting binned means of prioritization against the Labour lead, with raw observations displayed underneath.

Figure 7 plots these estimates across all three windows, for both the interaction and quadratic specifications. The estimates are stable but less precise as the window narrows from $[-0.20, 0.35]$ to $[-0.10, 0.25]$ under either specification. The F-statistic falls as the window narrows, reflecting the smaller effective sample at narrower bandwidths. At the narrowest window, closest to the MSE-optimal bandwidth, the F-statistic at the swing threshold falls to 9 under the interaction specification, below the F-statistic threshold of 10. The $[-0.20, 0.35]$ window carries the largest sample of the three, 374 against 312 and 229, and the point estimates vary by no more than 0.06 across the three windows.

5.2 Reduced form

I next examine whether crossing the thresholds is itself associated with a shift in 2024 general election outcomes, before combining this with the first stage in the next section to recover the Wald estimate.

Figure 8 plots the reduced-form relationship between the Labour lead and vote share or victory margin for Labour and the Conservatives, with a 95% confidence band around each segment and the raw data shown alongside it.

At the stronghold threshold, where the Labour campaign discontinuously stops, the Conservative vote share and victory margin rise significantly and consistently across all three




Figure 6: First-stage discontinuity in prioritization.




Figure 7: First-stage point estimates and confidence intervals across bandwidths.



Figure 8: Reduced-form relationships.

windows and across specifications, by 3.1 to 4.3 percentage points in vote share and 4.2 to 5.8 points in victory margin. Labour's victory margin falls by a comparable amount, a significant decline of 3.5 to 4.5 points in every window. Labour's vote share declines too, but more modestly and without reaching significance, by 1.2 to 1.6 points. Because crossing the stronghold threshold entails exiting treatment, the negative reduced-form coefficient on Labour's margin combines with the negative first stage to imply a positive effect of treatment on Labour's margin, taken up formally in Section 5.3. The positive reduced-form coefficients for the Conservatives, combined with the same negative first stage, implies a negative effect of treatment on the Conservatives' performance.

Figure 9 reports reduced-form estimates of party vote shares and turnout at the preferred bandwidth $[-0.20, 0.35]$.

At the swing threshold, where Labour begins campaigning, neither party's vote share moves detectably. The Conservative margin falls by 4.0 points at the main window, an estimate that is unstable across bandwidths, weakening to a 1.6-point decline at $[-0.15, 0.30]$ and reversing to a statistically insignificant 2.2-point rise at $[-0.10, 0.25]$.

The positive effect for Labour and negative for the Conservatives at the stronghold threshold, absent any significant turnout effect, may suggest direct movement between the two




Figure 9: Reduced-form effects on party vote shares and turnout.

parties and the success of Labour’s explicit strategy of persuading voters who once voted Conservative. Drawing solely on aggregate data at the constituency level risks an ecological fallacy. Influential actors such as the think tank Labour Together drew this conclusion from aggregate results, crediting the party’s victory in significant part to voters switching directly from the Conservatives. Existing survey evidence complicates that inference. Section 5.6 couples individual-level data from the British Election Study with the fuzzy RD design to test this directly.

5.3 Wald estimates

The Wald estimates are the ratio of the reduced-form coefficient to the first-stage coefficient at each cutoff, $\tau_j = \gamma_j^{RF} / \gamma_j^{FS}$ (equation (3), Section 4.3).

Figure 10 panels (a) and (b) report the preferred Wald estimates at bandwidth $[-0.20, 0.35]$, for victory margin and vote share outcomes, at both thresholds.

At the stronghold threshold, the Wald estimate of Labour’s victory margin is positive and stable in magnitude across all three bandwidths, ranging from 6.8 to 9.0 points. Only the main window reaches conventional significance. As noted previously, while the reduced-form coefficient stays significant throughout, the first-stage standard error grows and its F-statistic falls as the window narrows, mechanically inflating the Wald standard error. The point estimate for Labour’s vote share ranges from 2.3 to 3.1 points across windows, but is imprecise throughout.

The Conservative victory margin falls significantly across windows, by 8.2 to 11.5 points,




Figure 10: Wald estimates on vote shares and victory margins. Panel (a): victory margin. Panel (b): vote share.

and the Conservative vote share falls by 6.0 to 8.5 points, significant in the main and narrowest windows and borderline in the middle one. Reform’s vote share estimates are positive and consistent across all three windows, from 3.5 to 5.2 points, but none of the estimates reaches significance. Under the CCT robust bias-corrected estimator, the four coefficients are consistent with the main specification, though larger and less precise given the smaller sample within the MSE-optimal bandwidth.

At the stronghold threshold, Labour’s victory margin increases while the Conservative margin declines, in line with Labour’s stated strategy of targeting former Conservative voters. However, Reform’s positive point estimate at the same threshold suggests the aggregate pattern need not reflect that intended mechanism.

Around the stronghold threshold, the average Labour victory margin is larger than the estimated local effect. However, within 2.5 percentage points of the threshold, 8 of 32 constituencies were decided by less than the estimated effect on Labour’s victory margin, rising to 15 of 70 within a five-point window, implying some additional seats may have been won due to the campaign.

At the swing threshold, none of the four Wald estimates reaches significance in any window. Both Labour’s and the Conservative’s vote share estimates are unstable, with Labour’s estimates changing sign as the window narrows. Because victory margins are constructed from these same fluctuating vote shares, they provide little additional information.

Overall, I find a significant aggregate electoral gain for Labour at the stronghold threshold but no detectable corresponding gain at the swing threshold. If effect size declines with competitiveness, we may be overstating the count of 18 seats where an effect of average size would have been decisive. Sections 5.4 and 5.6 take up why the campaign seems to affect voting behavior at the stronghold threshold but not at the swing one. Both point to the roles of local competitive landscape and bloc dynamics, one using aggregate data and the other individual voting behavior.

5.4 Heterogeneity by contest type

All estimates so far are based on constituencies where the pre-campaign MRP prediction identified a major party, Labour, the Conservatives, Reform, or the Liberal Democrats, as a plausible contender. This section restricts the sample to seats where Labour and the Conservatives had already crystallized as the top two choices by 2019. These seats, which comprise most of the country, are also where a tactical effect should be most visible. Because Labour is the only alternative to defeating the Conservatives in these seats, the presence of strategic behavior should increase voting for Labour.

I estimate these results by running the same first-stage, reduced-form, and Wald sequence on this subsample alone. A specification that adds contest type as an interaction to the main specification cannot be estimated with the available sample. Splitting the sample instead costs some precision but is estimable.

At the stronghold threshold, the pattern closely mirrors the main results. Labour’s victory margin rises by 8.3 points, compared to 6.8 points in the full sample. Both estimates are significant in the main window. The Conservative margin falls significantly by 8.1 points, compared to 8.2 points in the full sample. Labour’s vote-share estimate is not significant in either sample, but it is larger at 3.6 points in this subsample relative to 3.1 points in the full sample. The Conservative vote-share estimate in this subsample falls by 5.1 points, close to the 6.1-point decline in the full sample, but falls just short of conventional significance. The loss of precision may reflect the smaller sample in the subsample. The Liberal Democrat and Reform vote shares show estimates similar to the main sample and neither are significant.

At the swing threshold, neither Labour’s margin nor its vote share is significant in any window, matching the null result in the full sample. However, Labour’s vote share estimate is consistently positive in this subsample. The Conservative vote share rises significantly by 9.2 points in the narrowest window. The Liberal Democrat margin falls significantly across all three windows, 10.7 to 12.3 points. Reform’s margin also falls, significant in the two narrower windows, mirroring the full-sample pattern.

Labour’s own vote share is similar at both thresholds, albeit small and not significant at either. The data cannot reject equality between the two estimates.

One reading of the Conservative and Reform results is that the campaign triggers consolidation behind the Conservatives in the most contested seats. In close Labour–Conservative races, Conservative-leaning voters who might otherwise vote Reform in protest against the government may instead back the Conservative candidate. While aggregate data cannot directly confirm that these individual voters are the ones switching, Section 5.6 adds more detail to this hypothesis using individual-level data.

Within the Labour–Conservative subsample, seats across the two thresholds differ in the 2019 Conservative vote share. To account for this imbalance, every specification in this section adds the Conservative vote share in 2019 directly, alongside the same controls used throughout the paper, including whether the seat was already prioritized before the short campaign.




Figure 11: Wald estimates in the restricted Labour-Conservative sample.

5.5 Robustness

Appendix Table A1 summarizes how the Wald estimates for the paper’s two headline outcomes, Labour and Conservative victory margins, respond to re-estimating within the union of the CCT-optimal bandwidths, replacing the interaction specification with a quadratic, and replacing the global polynomial with the CCT robust bias-corrected local polynomial estimator.

Covariate balance at the threshold is addressed in Section 4.4.2. The placebo tests below augment these balance checks by testing for discontinuities in a fuller set of 2019 election results.

An alternative bandwidth, the union of the CCT-optimal bandwidths at each cutoff, and an alternative estimator, the CCT robust bias-corrected local polynomial approach, both leave the stronghold-threshold results the same sign as the main specification throughout. Some estimates lose significance, which may be due to the smaller effective sample size. Point estimates also run larger, consistent with the weaker first stage that inflates standard errors encountered earlier.

An alternative functional form, with a quadratic specification in place of the interaction, is consistent with the main results at the stronghold threshold.

Placebo outcomes, 2017 and 2019 vote shares that predate the 2024 campaign, show no consistent discontinuity at the stronghold threshold, across window, sample, and year. One exception is the Conservatives’ 2019 vote share within Labour–Conservative contests at the swing threshold. To address this imbalance, I control for it directly in the Labour–Conservative specifications in Section 5.4.

Placebo thresholds test whether prioritization or the outcomes jump anywhere else along the running variable, apart from the two real cutoffs. Prioritization jumps significantly at both real cutoffs and not at placebo ones, in both samples ($p = .010$ throughout). Labour’s margin is significant at the stronghold threshold, in both the full sample ($p = .019$) and, more strongly, within Labour–Conservative contests ($p = .010$). The Conservative margin and vote shares move in the expected direction but fall short of significance. At the swing threshold, none of the four outcomes is significant in either sample. The placebo threshold specifications also utilize a smaller sample than the main result.

Overall, at the stronghold threshold, results are broadly consistent with the main results in sign and magnitude. Where significance is lost, it may reflect the smaller effective sample. The estimates at the swing threshold are less stable across specifications. Within Labour–Conservative contests specifically, both Labour’s and Conservative’s vote shares are positive under the alternative bandwidth, the CCT estimator, the quadratic specification, and placebo thresholds, and more stable than in the full sample. Conservative’s vote share

is also significant under the narrowest bandwidth. Narrowing the window shifts the sample increasingly toward Labour–Conservative contests, resulting in a smaller but more homogeneous sample. However, the vote share results for Liberal Democrats and Reform are not as consistent across specifications and therefore the margin of victory results remain noisy.

5.6 Corroborating evidence from the British Election Study on vote switching

The constituency-level results cannot distinguish between two potential mechanisms underlying the effect at the stronghold threshold. Changes in party vote shares may reflect voters switching parties, whether through persuasion or strategic behavior, or differential turnout. Separating these channels is a long-standing question in the study of effects of campaigns (Gerber and Green, 2000; Kalla and Broockman, 2018; Hill et al., 2021; Bonica and Grumbach, 2026). Individual-level panel data from the British Election Study allow me to track respondents’ electoral behavior between 2019 and 2024 and examine the individual-level changes underlying the aggregate effects. Combined with the constituency-level estimates, these data provide a link between individual electoral behavior and the constituency-level outcomes that determine representation.

Because the panel records respondents’ reported votes in both 2019 and 2024, it allows me to track movement between parties and relate these individual-level changes to the aggregate effects. The analysis includes 12,341 respondents in the full sample, 10,361 of whom live in constituencies with Labour–Conservative contests.

The BES results in the full sample are broadly consistent with the constituency-level estimates, with Conservative support declining at the stronghold threshold. At the swing threshold, the decline in Labour support is more pronounced in the BES data. The panel data also allow us to examine which voters account for these changes and whether they reflect switching between parties or differential turnout.

Among 2019 Conservative voters across all contests, the stronghold threshold produces a clear shift from the Conservatives to Reform, matching the overall result in the constituency-level data. The probability of voting Conservative in 2024 falls by 11.3 percentage points, while support for Reform rises by 10.6 points, with both effects statistically significant. At the swing threshold, Conservative support instead rises by 4.4 points, although the estimate is imprecise and there is no clear corresponding decline for any single party.

Restricting 2019 Conservative voters to Labour–Conservative contests clarifies the pattern at the swing threshold. Conservative support rises by 8.0 points (although not significantly), while Liberal Democrat support declines 7.3 points. At the stronghold threshold,




Figure 12: RDD estimates of vote choice using British Election Study data. Standard errors are clustered at the constituency level.

Conservative support again falls substantially while Reform support rises, with both effects statistically significant.

Among 2019 Liberal Democrat and Green voters, the full sample indicates movement from the Liberal Democrats toward Labour at the stronghold threshold, but little systematic switching at the swing threshold. In Labour–Conservative contests, these voters instead move from the Liberal Democrats toward Labour at both thresholds. At the swing threshold, Conservative support instead rises by 4.4 points, although the estimate is imprecise and there is no clear corresponding decline for any single party.

Additionally, 2019 Labour voters move significantly towards Liberal Democrats at the swing threshold in the full sample, but not in the Labour–Conservative sample. No other group of 2019 voters shows significant or consistent movement towards a party.

Taken together, the results point to strategic switching within ideological blocs, with both the direction of switching and the groups that respond depending on the electoral context. At the stronghold threshold, some former Conservative voters move toward Reform when the Conservatives are less viable, while former Liberal Democrat and Green voters consolidate behind Labour. Once the Conservatives become competitive in Labour–Conservative races, some former Conservative voters return to the party while progressive voters continue to consolidate behind Labour. The weaker aggregate return to Labour’s campaigning at the swing threshold is consistent with two competing responses to the campaign. Campaigning




Figure 13: 2019 Conservatives' post-election self-reported vote choice. Panel (a): all contests. Panel (b): Labour-Conservative contests.




Figure 14: 2019 Liberal Democrats' or Greens' post-election self-reported vote choice. Panel (a): all contests. Panel (b): Labour-Conservative contests.

consolidates support on the left, but also prompts Conservative-leaning voters to consolidate behind their own party, offsetting some of Labour’s gains.

6 Alternative explanations

6.1 Conservatives’ actions do not counteract Labour’s campaign

One possible explanation for the heterogeneous effects across thresholds is that they reflect differences in the Conservatives’ response to Labour’s campaign rather than differences in voters’ behavior. Section 4.4.3 provides little evidence for this mechanism, as neither Conservative Facebook spending nor campaign visits change discontinuously at either prioritization threshold (Figure 5).

I additionally re-estimate the main specifications controlling directly for both measures of Conservative campaign activity. Figure 15 reports the vote share results for the full sample and for Labour–Conservative contests. The estimates remain similar to the baseline results across bandwidths, although some are estimated less precisely. At the stronghold threshold, Labour’s gains and the corresponding decline in Conservative support remain, while the swing thresholds still display the same indeterminate patterns. Much as in the specifications that do not control for the Conservatives’ actions, restricting the sample to Labour–Conservative contests stabilizes the swing thresholds coefficients, with positive estimates for Labour and Conservative vote share and negative estimates for Reform and the Liberal Democrats. The estimates remain insignificant.

These results suggest that the difference between the two thresholds is not explained by the Conservatives offsetting Labour’s campaign around the swing threshold or potentially withdrawing from constituencies around the stronghold threshold.

6.2 Composition of voters across thresholds

One possible alternative explanation for the differing effects of the campaign at the two thresholds is that the constituencies around each contain different types of voters. For example, voters’ preferences for redistribution tend to rise with Labour’s predicted lead (see Appendix Figure A6). The larger effects at the stronghold threshold could thus reflect different preferences of voters living around the two thresholds rather than differences in how similar voters respond to the campaign.

I examine this possibility using individual-level data from the British Election Study. For each of four pre-campaign political preferences, I divide respondents into those who support and those who oppose redistribution, taxation and spending, immigration, and European




Figure 15: Constituency-level Wald estimates, controlling for the Conservatives' campaign activity. Panel (a): full sample. Panel (b): Labour–Conservative contests.

integration (corresponding to views on Brexit). I then estimate the campaign effect at each threshold separately for each group. Figure 16 shows the results for the two groups defined by redistribution preferences, while the results for the remaining groups are reported in Appendix Figure A7.




Figure 16: RDD based on BES data for voters who, in panel (a), oppose redistribution and, in panel (b), support redistribution.

Among respondents who support redistribution, for example, the effect on Labour voting shifts significantly from -20 to $+11$ percentage points across the two thresholds. More generally, I can reject equality of the campaign effects for five of the eight preference groups

on at least one electoral outcome. Even where differences are imprecisely estimated, the pattern of party choice generally looks substantively different across thresholds, with estimates frequently changing sign across thresholds.

6.3 Social networks

Another possible explanation for the difference in campaign effects across the two thresholds is that campaign activity spreads through voters' social networks. Campaign contact may have larger indirect effects where voters discuss politics more or have more opportunities to do so with politically sympathetic people. A campaign message received by one voter could then reach or be reinforced by a larger number of potentially receptive voters.

Figure 17 provides little evidence that social networks could explain the difference between the two thresholds. The share of respondents who discuss politics at least once a week is broadly stable across Labour's predicted margin, with no visible change at either prioritization threshold. The same is true when respondents are separated by their party identification (Figure 17, panel b). Among Labour, Conservative, and Liberal Democrat identifiers, there is no clear increase in the frequency of political discussion as Labour's predicted lead increases, while the estimates for Reform identifiers are noisier due to the small sample.

I also examine the political loyalties of the discussants in Appendix Figures A8 and A9. While the previous analysis shows how frequently respondents identifying with different parties discuss politics with anyone, it does not identify the political views of the people with whom they discuss it. Here I consider the share of respondents' political discussants who support the same party as the respondent. While this measure is noisier due to the much smaller effective sample, the share of co-partisan discussants remains broadly stable across Labour's predicted margin, both overall and within parties, with no clear change at either threshold.

Finally, I estimate campaign effects separately by respondents' frequency of political discussion and the partisan composition of their discussion networks. If political discussion generates the observed differences across thresholds, campaign effects should vary with the amount and nature of political conversations and become more similar across the two thresholds when comparing respondents with similar discussion levels. Appendix Figure A10 reports the results by frequency of political discussion and provides little evidence of either pattern. The estimated effects on Labour and Conservative voting do not become systematically more pronounced as the frequency of political discussion increases, and differences between the two thresholds remain among respondents with similar levels of political discussion. Appendix




Figure 17: Political discussion intensity by Labour lead. Panel (a): full sample. Panel (b): split by party ID.

Figure A11 repeats the analysis according to the share of co-partisans in respondents' discussion networks and produces a similar pattern. Together with the absence of meaningful differences in the prevalence of discussion across Labour's predicted margin, these results provide little evidence that political discussion networks explain the larger campaign effect at the stronghold threshold.

7 Further evidence on mechanisms and downstream effects

7.1 Party and government evaluations

I next examine whether campaign prioritization changed voters' evaluations of the parties and the incumbent government. Figure 18 reports effects on liking Labour and the Conservatives, retrospective evaluations of the economy, approval of the UK government, and expectations that the Conservatives would achieve four major policy goals. At the swing threshold, there is little evidence that campaigning changed any of these evaluations. At the stronghold threshold, campaigning significantly reduced liking the Conservatives and produced more negative assessments of the economy and the government. Voters also became substantially less likely to expect the Conservatives to achieve their policy goals. Liking Labour, on the other hand, is essentially unchanged. The same pattern is visible when the sample is restricted to Labour–Conservative contests. Thus, where campaigning produced the larger electoral effects, it also appears to have made voters more negative about the Conservatives and their performance, without making them more positive about Labour.

These results do not establish that changes in political evaluations caused changes in vote choice, but they suggest a potential difference in how campaigning operated at the two thresholds. This pattern is consistent with the voting behavior documented above. Around the swing threshold, where electoral competition creates stronger incentives for strategic voting in Labour–Conservative contests, campaigning changes voters' underlying evaluations relatively little and appears to produce more strategic voting behavior. At the stronghold threshold, voters respond to the campaign with expressive behavior, while campaigning also makes voters more negative about the Conservatives and their performance.

Additional evidence on voters' perceptions of electoral viability points in the same direction (Appendix Figure A12). The estimates are imprecise, but at the swing threshold the belief that votes for smaller parties are wasted moves upward, while at the stronghold threshold reported strategic voting moves downward.




Figure 18: The effects of the campaign on party evaluations.

7.2 Effects on behavior of Members of Parliament

Campaign prioritization also appears to affect the subsequent behavior of Labour MPs. I examine whether Labour MPs publicly called on Keir Starmer to resign during the leadership crisis starting in May 2026. Calls for his resignation followed a series of unpopular government decisions, including cuts to winter fuel payments and disability benefits, criticism within the party of the government’s rhetoric on immigration, and revelations that Starmer’s choice for ambassador to the United States had close ties to Jeffrey Epstein.

Figure 19 shows the raw relationship between Labour’s MRP-predicted lead and MPs calling for Starmer to resign. Below the stronghold threshold, the share rises with Labour’s predicted lead, reaching around 40 percent among MPs in constituencies with a predicted lead of 10–15 percentage points. Immediately above the threshold, it falls to around 15 percent. There is no comparable discontinuity at the swing threshold.




Figure 19: Percentage of MPs calling for Starmer to resign, by Labour lead.

The RD estimates confirm this pattern. At the stronghold threshold, prioritization increases the probability that an MP calls for Starmer to resign by approximately 60 percentage points in the interacted specification. The estimate is statistically significant and remains similar across bandwidths, although it loses significance at the narrowest bandwidth. There is no significant effect at the swing threshold.

This result is notable because prioritized seats received the greatest support and resources from the central campaign. If receiving this additional support generated loyalty to the party




Figure 20: Wald estimates of the effect of the campaign on calling for Starmer to resign.

leadership, prioritization should have made MPs less likely to call for Starmer’s resignation. The opposite pattern could have arisen because the campaign changed the electorate to which MPs subsequently had to respond. The results in earlier sections indicate that the campaign drew voters from other progressive parties toward Labour, particularly where Labour was already comfortably ahead. This change in Labour’s electoral coalition may have increased MPs’ incentives to oppose an unpopular party leadership and remain accountable to their voters.

8 Conclusion

This paper provides causal evidence on the electoral effects of the full set of campaign resources deployed by a major political party. Exploiting discontinuities in Labour’s allocation of resources during the 2024 UK general election, I find that campaigning can substantially change electoral outcomes, but that its returns depend on the competitive environment in which it takes place. Campaign prioritization increases Labour’s victory margin by 6.8 percentage points where the party is already comfortably ahead, while producing no comparable aggregate gain in the closest races. The election took place amid an unusually broad

anti-incumbent shift across developed democracies, which may have made voters particularly willing to move away from the Conservatives (Burn-Murdoch, 2024). The contrast within the same election nevertheless shows that the same campaign can generate strikingly different electoral returns.

More broadly, the results show that the return to campaigning depends on the strategic responses it induces among voters as well as on the support it yields for the campaigning party itself. In Labour–Conservative contests, former Liberal Democrat and Green voters consolidate behind Labour at both thresholds, while former Conservative voters move toward Reform when Labour is comfortably ahead but consolidate behind the Conservatives in close contests. This suggests that campaign effects are endogenous to the competitive environment in which the campaign takes place. The individual-level results also reveal that the aggregate decline in Conservative support does not primarily reflect direct switching from the Conservatives to Labour, despite Labour’s explicit effort to win former Conservative voters.

These results also have implications for campaign strategy. Parties have strong incentives to concentrate resources in close races and on voters they consider persuadable, but electoral competitiveness can itself shape how those voters respond to campaigning. The electoral return to campaigning in the closest races may therefore be limited by countervailing responses among voters aligned with the opposing party.

Finally, campaign allocation has consequences beyond election day. In constituencies where Labour was comfortably ahead, prioritization substantially increases the probability that the Labour MP subsequently calls for Prime Minister Keir Starmer to resign during the 2026 leadership crisis. This may reflect changes in the electoral coalition facing MPs after the election. Campaigns can affect not only voters and electoral outcomes, but also the subsequent political behavior of elected representatives.

References

- Aggarwal, M., Allen, J., Coppock, A., Frankowski, D., Messing, S., Zhang, K., Barnes, J., Beasley, A., Hantman, H., and Zheng, S. (2023). A 2 million-person, campaign-wide field experiment shows how digital advertising affects voter turnout. *Nature Human Behaviour*, 7(3):332–341.
- Alt, J., Bueno de Mesquita, E., and Rose, S. (2011). Disentangling accountability and competence in elections: Evidence from u.s. term limits. *The Journal of Politics*, 73(1):171–186.
- Alvarez, R. M., Hopkins, A., and Sinclair, B. (2010). Mobilizing pasadena democrats: Measuring the effects of partisan campaign contacts. *The Journal of Politics*, 72(1):31–44.
- Ansolabehere, S. and Iyengar, S. (1995). *Going Negative: How Political Advertisements Shrink and Polarize the Electorate*. Free Press, New York.
- Ansolabehere, S., Iyengar, S., Simon, A., and Valentino, N. (1994). Does attack advertising demobilize the electorate? *American Political Science Review*, 88(4):829–838.
- Arceneaux, K. (2007). I’m asking for your support: The effects of personally delivered campaign messages on voting decisions and opinion formation. *Quarterly Journal of Political Science*, 2(1):43–65.
- Arceneaux, K. and Nickerson, D. W. (2009). Who is mobilized to vote? a re-analysis of 11 field experiments. *American Journal of Political Science*, 53(1):1–16.
- Bailey, M. A., Hopkins, D. J., and Rogers, T. (2016). Unresponsive and unpersuaded: The unintended consequences of a voter persuasion effort. *Political Behavior*, 38(3):713–746.
- Besley, T. and Case, A. (1995). Does electoral accountability affect economic policy choices? evidence from gubernatorial term limits. *Quarterly Journal of Economics*, 110(3):769–798.
- Black, A. (2024). General election result, campaign strategy, conference – Labour NEC report. LabourList. Black’s summary, as an NEC member, of the July 23, 2024 National Executive Committee meeting, at which Ellie Reeves and General Secretary David Evans briefed the NEC that activists had “spoke to 3.1 million voters between May 22nd and July 4th alone.”
- Bonica, A. and Grumbach, J. M. (2026). The electoral effects of candidate ideology in the Trump era. Working paper, dated August 3, 2026.

- Bowler, S. (2010). Private members’ bills in the uk parliament: Is there an ‘electoral connection’? *The Journal of Legislative Studies*, 16(4):476–494.
- Broockman, D. E. and Kalla, J. L. (2023). When and why are campaigns’ persuasive effects small? evidence from the 2020 u.s. presidential election. *American Journal of Political Science*, 67(4):833–849. Online-first in 2022; print/citable year is 2023 – see bib_notes.md.
- Burn-Murdoch, J. (2024). Democrats join 2024’s graveyard of incumbents. Financial Times.
- Calonico, S., Cattaneo, M. D., and Titiunik, R. (2014). Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica*, 82(6):2295–2326.
- Cantoni, E. and Pons, V. (2021). Strict id laws don’t stop voters: Evidence from a u.s. nationwide panel, 2008–2018. *Quarterly Journal of Economics*, 136(4):2615–2660. Circulated earlier as NBER Working Paper No. 25522 (2019) – see bib_notes.md if you meant to cite the working-paper version.
- Cattaneo, M. D., Jansson, M., and Ma, X. (2018). Manipulation testing based on density discontinuity. *The Stata Journal*, 18(1):234–261.
- Cattaneo, M. D., Titiunik, R., Vázquez-Bare, G., and Keele, L. (2016). Interpreting regression discontinuity designs with multiple cutoffs. *The Journal of Politics*, 78(4):1229–1248.
- Coppock, A., Green, D. P., and Porter, E. (2022). Does digital advertising affect vote choice? evidence from a randomized field experiment. *Research & Politics*, 9(1). Article 20531680221076901.
- Coppock, A., Hill, S. J., and Vavreck, L. (2020). The small effects of political advertising are small regardless of context, message, sender, or receiver: Evidence from 59 real-time randomized experiments. *Science Advances*, 6(36):eabc4046.
- Cracknell, R. and Baker, C. (2024). General election 2024: Results and analysis. Research Briefing CBP-10009, House of Commons Library, London.
- Cruz, C., Labonne, J., and Trebbi, F. (2026). Campaigning to a polarised electorate: Emotions and information in real election campaigns. *The Economic Journal*. Forthcoming; advance article, DOI: 10.1093/ej/ueag058. Circulated earlier as NBER Working Paper No. 32112 under the title “Campaigning Against Populism.”.
- Dixit, A. and Londregan, J. (1996). The determinants of success of special interests in redistributive politics. *The Journal of Politics*, 58(4):1132–1155.

- Downs, A. (1957). *An Economic Theory of Democracy*. Harper & Row, New York.
- Electoral Commission (2025). General election spending hits record high. Press release. Reports £94.5 million in total spending by parties and non-party campaigners at the 2024 UK general election.
- Enos, R. D. and Fowler, A. (2018). Aggregate effects of large-scale campaigns on voter turnout. *Political Science Research and Methods*, 6(4):733–751.
- Enríquez, J. R., Larreguy, H., Marshall, J., and Simpser, A. (2024). Mass political information on social media: Facebook ads, electorate saturation, and electoral accountability in mexico. *Journal of the European Economic Association*, 22(4):1678–1722.
- Esteban-Casanelles, T. (2024). The electoral effects of street advertising: Quasi-experimental evidence from barcelona. Working paper, King’s College London. *Confirm exact title/year against the version you intend to cite – see bib_notes.md*.
- Ferraz, C. and Finan, F. (2011). Electoral accountability and corruption: Evidence from the audits of local governments. *American Economic Review*, 101(4):1274–1311.
- Ford, R., Bale, T., Jennings, W., and Surridge, P. (2025). *The British General Election of 2024*. Palgrave Macmillan, Cham.
- Fourinaies, A. and Hall, A. B. (2022). How do electoral incentives affect legislator behavior? evidence from u.s. state legislatures. *American Political Science Review*, 116(2):662–676.
- Gerber, A. S., Gimpel, J. G., Green, D. P., and Shaw, D. R. (2011). How large and long-lasting are the persuasive effects of televised campaign ads? results from a randomized field experiment. *American Political Science Review*, 105(1):135–150.
- Gerber, A. S. and Green, D. P. (2000). The effects of canvassing, telephone calls, and direct mail on voter turnout: A field experiment. *American Political Science Review*, 94(3):653–663.
- Green, D. P. and Gerber, A. S. (2004). *Get Out the Vote: How to Increase Voter Turnout*. Brookings Institution Press, Washington, DC.
- Hager, A. (2019). Do online ads influence vote choice? *Political Communication*, 36(3):376–393.
- Hahn, J., Todd, P., and van der Klaauw, W. (2001). Identification and estimation of treatment effects with a regression-discontinuity design. *Econometrica*, 69(1):201–209.

- Hewitt, L., Broockman, D., Coppock, A., Tappin, B. M., Slezak, J., Coffman, V., Lubin, N., and Hamidian, M. (2024). How experiments help campaigns persuade voters: Evidence from a large archive of campaigns’ own experiments. *American Political Science Review*, 118(4):2021–2039.
- Hill, S. J., Hopkins, D. J., and Huber, G. A. (2021). Not by turnout alone: Measuring the sources of electoral change, 2012 to 2016. *Science Advances*, 7(17):eabe3272. Topic is composition- vs. conversion-based sources of electoral change – see bib_notes.md.
- Imbens, G. W. and Angrist, J. D. (1994). Identification and estimation of local average treatment effects. *Econometrica*, 62(2):467–475.
- Imbens, G. W. and Lemieux, T. (2008). Regression discontinuity designs: A guide to practice. *Journal of Econometrics*, 142(2):615–635.
- Kalla, J. L. and Broockman, D. E. (2018). The minimal persuasive effects of campaign contact in general elections: Evidence from 49 field experiments. *American Political Science Review*, 112(1):148–166.
- Kellermann, M. (2013). Sponsoring early day motions in the british house of commons as a response to electoral vulnerability. *Political Science Research and Methods*, 1(2):263–280.
- Kellermann, M. (2016). Electoral vulnerability, constituency focus, and parliamentary questions in the house of commons. *The British Journal of Politics and International Relations*, 18(1):90–106.
- Kendall, C., Nannicini, T., and Trebbi, F. (2015). How do voters respond to information? evidence from a randomized campaign. *American Economic Review*, 105(1):322–353.
- Larreguy, H. A., Marshall, J., and Snyder, James M., J. (2018). Leveling the playing field: How campaign advertising can help non-dominant parties. *Journal of the European Economic Association*, 16(6):1812–1849.
- Lee, D. S. and Lemieux, T. (2010). Regression discontinuity designs in economics. *Journal of Economic Literature*, 48(2):281–355.
- List, J. A. and Sturm, D. M. (2006). How elections matter: Theory and evidence from environmental policy. *The Quarterly Journal of Economics*, 121(4):1249–1281.
- Miori, M. and Green, J. (2025a). The most disproportionate uk election: How the labour party doubled its seat share with a 1.6-point increase in vote share in 2024. *The Political Quarterly*, 96(1):37–64.

- Miori, M. and Green, J. (2025b). Party bloc equilibrium: Britain’s quiet electoral earthquake. *Political Insight*, 16(4):4–8.
- Nickerson, D. W. (2008). Is voting contagious? evidence from two field experiments. *American Political Science Review*, 102(1):49–57.
- Núñez, L. (2025). Encouraging loyalty and defection: The impact of campaigns on tactical voting in britain. *American Journal of Political Science*, 69(3):981–994.
- OpenSecrets (2024). Total 2024 election spending projected to exceed previous record. OpenSecrets News. Projects at least \$15.9 billion in total federal election spending for the 2024 cycle.
- Pons, V. (2018). Will a five-minute discussion change your mind? a countrywide experiment on voter choice in france. *American Economic Review*, 108(6):1322–1363.
- Spenkuch, J. L. and Toniatti, D. (2018). Political advertising and election results. *Quarterly Journal of Economics*, 133(4):1981–2036.
- Staiger, D. and Stock, J. H. (1997). Instrumental variables regression with weak instruments. *Econometrica*, 65(3):557–586.

Appendix




Figure A1: Conservative vote share against Labour's MRP-predicted margin, showing Scottish constituencies clustering at lower Conservative support (referenced in Section 4.1).

Table A1: Robustness of the Wald estimates to bandwidth, functional form, and estimator

	Labour victory margin (1)	Conservative victory margin (2)
Panel A: Swing threshold		
Main specification ($N = 374$)	-8.22 (5.17)	-10.65** (5.06)
Alternative bandwidth ($N = 222$)	-2.89 (5.07)	4.28 (5.02)
Quadratic ($N = 374$)	-9.30* (5.40)	-11.30** (5.02)
CCT robust ($N_h \approx 121$ – 124)	-17.23 (22.53)	16.22 (18.39)
Panel B: Stronghold threshold		
Main specification ($N = 374$)	7.33* (3.75)	-8.15** (3.70)
Alternative bandwidth ($N = 222$)	10.46 (6.63)	-11.45 (7.11)
Quadratic ($N = 374$)	9.16** (3.59)	-5.87* (3.22)
CCT robust ($N_h = 91$)	51.20 (42.32)	-49.36 (41.53)

Notes: Wald (fuzzy RD) estimates in percentage points. HC1 standard errors in parentheses, except CCT robust rows, which report the robust bias-corrected standard error of [Calonico et al. \(2014\)](#). “Main specification” and “Quadratic” use the $[-0.20, 0.35]$ window. “Alternative bandwidth” re-estimates the main (interaction) specification within the union of the CCT-optimal bandwidths at each cutoff (h^* swing = 0.096, giving $[-0.096, 0.096]$; h^* stronghold = 0.084, giving $[0.066, 0.234]$; same window and N for both outcomes and both panels). CCT robust instead uses the CCT-optimal bandwidth with the robust bias-corrected local polynomial estimator, so the effective local sample N_h differs by column and threshold. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Reduced form point estimates for vote shares of Labour, Conservative, Liberal Democrat, Reform and turnout




Figure A2: Reduced-form effects on party vote shares and turnout across the full set of bandwidths.





Figure A3: Wald estimates on vote shares and victory margins across the full set of bandwidths, all races.





Figure A4: Wald estimates on vote shares and victory margins across the full set of bandwidths, restricted Labour–Conservative sample.




Figure A5: Wald estimates controlling for Conservative campaign activity across the full set of bandwidths. Top panel: all races. Bottom panel: Labour–Conservative contests.




Figure A6: Voter preferences for redistribution by Labour's predicted lead, all-contests sample.




Figure A7: RDD estimates by remaining pre-campaign preference groups (taxation and spending, immigration, European integration), all-contests sample. Panel (a): support raising taxes and spending. Panel (b): support cutting taxes and spending. Panel (c): support more immigration. Panel (d): support restricting immigration. Panel (e): pro-EU integration. Panel (f): pro-EU independence. Companion to Figure 16, which reports the redistribution dimension of the same eight preference groups.




Figure A8: Political loyalties of discussants, by respondent party ID: share of each respondent's political discussants who support the same party, all-contests sample. Dashed line marks a 50% co-partisan split.




Figure A9: Political loyalties of discussants (continued): the same co-partisan discussant share as Appendix Figure A8, pooled across all respondents rather than split by party ID.




Figure A10: Campaign effects by frequency of political discussion, all-contests sample. Panel (a): no political discussion (0 days/week). Panel (b): some discussion (1–3 days/week). Panel (c): frequent discussion (4–7 days/week).




Figure A11: Campaign effects by share of co-partisan discussants, all-contests sample. Panel (a): a majority (over 50%) of discussants are co-partisan. Panel (b): mixed discussion network (50% or fewer co-partisan).




Figure A12: Perceptions of electoral viability and strategic voting, all-contests sample: z-scored agreement that a vote for a small party is wasted, that it does not matter who is in power, and self-reported strategic voting (preferred another party but voted for the one with a chance of winning).